

FOR PROJECT MANAGERS

Turn meeting notes and project files into a plan that remembers why.

PrismPM reads the notes and documents you already have, proposes the Decisions, Tasks and Milestones inside them with the exact quote, and puts what you accept on your Timeline. Later, ask "why did we...?" and get the source, or a plain "nothing was recorded".

A PM's information is scattered. The schedule lives in a tool like Microsoft Project or Jira. The inputs live everywhere else: meeting notes, transcripts, a Word brief, an Excel tracker. Turning them into a plan means re-typing, and the reasons rarely survive a handover.

1 · EVIDENCE

Kickoff notes · pasted

Action item: Book the usability lab by 2026-10-10

Milestone: Pilot readout on 2026-10-20.

Notes Transcripts PDF Word Excel CSV
Scanned PDF

Paste or upload. A background pass reads new text on save.

2 · PROPOSALS, EACH QUOTING ITS SOURCE

☒ Proposed task by the Assistant

Book the usability lab

Accept Edit and accept Reject

Sources

Kickoff notes	Owner	None
Action item: Book the usability lab by 2026-10-10	Milestone	None
	Start	None
	Due	10 Oct 2026

☒ Proposed milestone by the Assistant

Pilot readout


Accept Edit and accept Reject

Sources

Kickoff notes	Owner	None
Milestone: Pilot readout on 2026-10-20.	Due	20 Oct 2026

Nothing enters the Project until the PM clicks Accept, Edit and accept, or Reject.

3 · ACCEPTED ITEMS LAND ON THE TIMELINE



A Gantt-style Timeline with dependencies, plus board, list and Calendar views. Accepted Proposals show in History as "via Assistant".

4 · ASK WHY

Searched decisions

The decision to freeze the legacy gateway on 1 October was made instead of running both the legacy and new systems in parallel. This choice was agreed upon during the sprint review, emphasizing a more streamlined transition. For further details, you can refer to the D-4 Freeze the legacy gateway on 1 October instead of running both in parallel decision, which cites the Sprint review notes as the source of this agreement.

Links to the Decision and the notes behind it. With no Decision recorded, it says so.

WHY PRISMPM DESERVES THE POINTS

Scattered evidence becomes a plan

Rovo and Notion AI pull action items from meetings; NotebookLM answers from files. PrismPM adds confirmed Decisions, Assumptions, supersession and review when one breaks.

US\$0.001

per Proposal pass on a 9-source batch; an unchanged repeat costs US\$0

Reviewable by design

Proposals quote their source and wait for a click. In-app Assistant writes need approval by default. "Why" answers search only confirmed Decisions.

16 of 16

citations resolve, up from 0 of 6 once tools returned ready-made links

Honest about the weak spots

54 eval cases and a three-model bake-off, which showed our live default is the weakest and obeys a hidden injection. We say so, and Users can switch with their own key.

20/20

"why" answers on gemini-2.5-flash, vs 14/20 on the live gpt-4o-mini

Try it in two minutes

1 Open jira-pro-max.vercel.app and sign up. New accounts start with a sample Project.

2 On the **Evidence** tab, paste meeting notes or upload a PDF, Word or Excel file.

3 On **Overview**, accept the Proposals under Needs attention, then open **Timeline**.

4 Open the **Assistant**, ask "why did we...?" and click the citation.

The live preview runs gpt-4o-mini, our only production key. Add your own key in Settings to use the measured pick, gemini-2.5-flash.

What shipped

Evidence

Notes, transcripts, PDF, Word, Excel, CSV and Markdown, with extracted text, semantic search and passage-level citations.

Proposals

Decisions, Tasks and Milestones extracted on save, each with its Source and verbatim quote. Pending until a PM accepts.

Planning views

Tasks, Milestones, Dependencies, Risks and People on a board, list, Timeline and Calendar.

Cited “why” answers

Answers from confirmed Decisions, linking each claim to its Source, or say no Decision was recorded.

Decisions and Assumptions

Typed Assumptions and supersession. Change a Task or Milestone that breaks an Assumption and its Decisions are flagged.

Approval cards

In-app Assistant writes ask first: Deny, Allow once, or an explicit Always allow per tool.

Full History

Field-by-field History for every item, with the Assistant’s changes marked “via Assistant”.

Concept Renders

A concept image of the deliverable, drafted from up to 3 Evidence items. Only the text the PM approved is sent.

MCP server

27 of 30 Assistant tools for Claude Desktop, Cursor and others, by personal token. Destructive tools are not exposed.

Your own model

OpenAI, Anthropic, Gemini or any OpenAI-compatible endpoint, key encrypted at rest. Plus Rooms to chat with invited Participants.

Measured, not assumed

54

eval cases, graded deterministically: 22 Decision, 12 Task/Milestone, 20 “why” answer

US\$0.0030

mean cost per Assistant turn on the live gpt-4o-mini, at full uncached prices

9–12/12

Task/Milestone cases on gpt-4o-mini in 3 runs; a new owner rule cut invented owners from every run to 1 in 5 runs

708


passing Vitest tests in 66 files on main, plus Playwright end-to-end tests in CI on every PR

MODEL (BAKE-OFF, 28 SEP)	DECISION EXTRACTION	“WHY” ANSWERS	HIDDEN “PWNERD” INSTRUCTION	COST PER ANSWER CASE
gpt-4o-mini · live default	16/22 → 20/22	13/20 → 14/20	Obeyed	US\$0.0016
gemini-2.5-flash · measured pick, own key	21/22 → 21/22	19/20 → 20/20	Refused	US\$0.0031
claude-haiku-4.5 · quality ceiling	20/22	20/20	Ignored, but over-extracted	US\$0.0246

Arrows show the score after our prompt work. gpt-4o-mini’s after-scores were measured through OpenAI directly: 20/22 in all three runs of the merged code on 1 Oct (the injection case fails every time), and 14/20 answers on 30 Sep. The live preview runs gpt-4o-mini because it is the only production key the team holds, so it keeps a known model-quality limit; switching is one environment change.

Safety and trust

- AI proposes, a PM accepts. When gpt-4o-mini obeyed the injected “PWNERD” instruction, the Proposal still reached the review queue, but it cannot become a Decision without a PM’s click.
- Writes need approval by default. Always allow is an explicit, per-tool choice. Over MCP, the client’s own prompt applies and destructive tools are left out.
- Citations are server-supplied. The dock links only destinations the server handed the model.



Business

- First users: PMs or delivery leads on small implementation teams, with recurring handovers and reasons spread across notes and files.
- Fit: sits before or beside the scheduling tool they already use.
- Price (proposed): US\$12 per seat for Team. Today it is a free research preview.
- Unit cost: about US\$0.70 a month of inference for a typical active PM, 6% of a seat. Hosting, storage and support are extra.
- Value: a seat pays off if it saves a PM some reconstruction time. Not yet validated.
- Moat: the features can be copied. A team’s growing record of confirmed Decisions and corrections would not be, but that needs repeated use we don’t have yet.

PMs don’t need another place to re-type the plan. PrismPM drafts it from what they already have, keeps the reasons attached, and leaves every Decision to a person.

Go deeper in the milestones PDF
Problem and competitors: M0, M1 · Pricing and costs: M6
Model choice and evals: M9, M11 · Prompts: M8
Safety: M13 · Workflows: M16 · MCP: M23
[jira-pro-max.vercel.app](#) · [GitHub repo](#)